Locality Causes Tractability? An Intervention Study on the Variable-Interaction Radius in Geometric Random Satisfiability

Jin Song
Independent Researcher
j.song.cs@outlook.com

September 2026
Compact Version — a full journal version is in preparation.

Abstract

A widespread intuition holds that real-world computational problems are easy because their "causes act locally": when variables interact only within small neighborhoods, constraint solvers should scale well. Observational evidence supports the intuition, but observation cannot separate locality from the many other regularities real instances carry. We translate the intuition into a falsifiable, interventional claim: holding the number of variables fixed, matching clause density relative to each radius's own threshold, and with clause-content statistics r-invariant by construction, varying *only* the radius r within which clause variables are drawn from a 2-D torus should causally change CDCL solving difficulty. Three findings. *(i)* The operational satisfiability boundary is not r-invariant: it lies below density 2.5 at $r = 0.06$ and near the classical $\approx$4.3 region by $r = 0.22$ (a displacement of at least 1.7 density units) — locality moves the sat/unsat boundary itself — before becoming unmeasurable at $r \geq 0.3$ under a 10^7-conflict budget. *(ii)* At density matched relative to each radius's own threshold, difficulty spans 3.4 orders of magnitude across r among decided instances (mean $\log_{10}$ conflicts $1.45 \to 4.86$ from $r = 0.06$ to 0.22), and ≥ 4.5 orders when budget-censored runs are counted at their lower bound; RM-ANOVA $F(8, 216) = 2808.5$, $p \approx 10^{-213}$; Jonckheere–Terpstra $p \leq 10^{-4}$ with FDR $q \leq 10^{-4}$; partial $\eta^2 \approx 0.99$. *(iii)* A degree-preserving swap that destroys spatial locality while keeping every literal's occurrence count exactly fixed raises difficulty by 16–68$\times$ (Wilcoxon $p \leq 7 \times 10^{-10}$) and flips the satisfiability status of 53 of the 82 resolved instance pairs; the ratio is a composite spanning those status changes (within the 29 status-stable pairs the elevation persists at $\times$11 and $\times$1.8 per offset arm, $p \leq 3.9 \times 10^{-3}$); post-hoc, exploratory causal mediation analysis is consistent with $\sim$94% of the radius effect flowing through community-structure restructuring (ACME 3.06, bootstrap CI $[2.84, 3.37]$). Locality is not merely correlated with tractability — under our controls the evidence is consistent with spatial locality contributing to the causal channel. All claims are scoped to $n = 400$, direct CNF encodings, and CDCL-family solvers; we discuss what this does and does not say about P versus NP. Code, seeds, data, preregistration, and the full audit trail are public.

1 Introduction

The conjecture and its observational gap. In popular discussions of P versus NP, one encounters a philosophical diagnosis: problems resist algorithmic solution when variable interactions are strongly correlated at long range; the computational world is tractable because causation is local. As stated, the claim is unfalsifiable — "locality", "correlation", and "cause" are unoperationalized. Yet the intuition has an empirically serious core. Industrial SAT instances, the reason conflict-driven clause-learning (CDCL) solvers [1] routinely decide formulas with millions of variables from

competition and application repositories [2], exhibit community structure, small backdoors, and heavy-tailed variable-incidence patterns. Observationally, structure and ease co-occur.

Observation, however, cannot license a causal reading. Instances that are easy and local are also many other things. The literature's own cautionary results make the point sharp. Zulkoski et al. [4], in the largest correlational study to date (∼7000 competition instances), found that *no single* structural parameter significantly predicts CDCL runtime across heterogeneous benchmark sets. The best six-feature regression reaches adjusted $R^2 \approx 0.31$ on application instances, and the size-plus-treewidth feature set reaches only 0.05. Mull, Fremont and Seshia [5] prove that any "polynomial-clique" community metric, including modularity, admits instances with excellent scores that remain NP-hard, and that pseudo-industrial community-attachment instances require exponentially long resolution proofs with high probability when communities are few. Structure, by itself, neither certifies nor explains tractability.

What an intervention adds. The remedy is standard in the empirical sciences and rare in SAT difficulty research: hold everything else fixed and turn one knob. We operationalize the locality conjecture as a dose–response experiment. A *locality-kernel generator* samples each clause's variables uniformly from an r-ball on a continuous 2-D unit torus and draws polarities from a standardized content distribution whose statistics are r-invariant by construction (σ-redundancy histograms identical across r in pilots). Literal-degree means equal 3α exactly, varying across r only through the threshold-matched density α, and the intervention arm's degree-preserving swap holds the degree profile exactly fixed. The knob r has a physical reading: it is the *spatial span of variable interaction* — the exact quantity the philosophical conjecture is about.

Contributions.

1. **To our knowledge, the first threshold-matched, preregistered intervention scan of locality radius.** Nine r levels × three preregistered offsets Δ above each radius's *own* operational threshold $\alpha_c(r)$ × 30 paired seeds × two solvers, with per-instance structural metrics (treewidth bracketed by min-fill upper and degeneracy lower bounds; modularity, community metrics, spectral gap). A fourth offset, Δ=+0.6, is included at 4 of the 9 radii; the r=0.06 anchors are the lattice-nearest S0 grid value 2.6 (the grid shows the true boundary lies below 2.5), which keeps the threshold-displacement claim a conservative lower bound and over-doses the r=0.06 scan cells by a bounded amount, on the order of 0.1–0.2 decades.
2. **A threshold curve** $\alpha_c(r)$ (832 solves at 10^7-conflict budget): the operational threshold is not r-invariant: the boundary lies below 2.5 at $r = 0.06$ and near ≈4.3 by $r = 0.22$, before difficulty saturates beyond measurable range at $r \geq 0.3$; the *measurability boundary itself* tracks r.
3. **A structure-destruction intervention (E3b).** A degree-preserving swap chain destroys spatial locality while keeping the literal degree sequence exactly fixed; paired Wilcoxon tests probe the causal contribution of the structure-destroying bundle.
4. **Topology × content decomposition**, with the content arm's operationalization registered openly as an amendment (AMEND-3) rather than shipped ill-defined.
5. **An honest negative-result protocol**: preregistration frozen before the main scan, numbered amendments, a prespecification-vs-implementation audit that caught one material bug before unblinding, and quantitative triggers under which we would have reported "density dominates".

Scope. All claims concern CDCL-family solvers on direct CNF encodings at $n = 400$ with conflict-count difficulty. We make no claim about P versus NP (§11). Machine-independence is obtained by using conflict counts, not wall-clock, as the primary measure.

2 Related Work

Structural measures and CDCL performance. Ansótegui et al. introduced the scale-free (power-law) variable-incidence model of industrial-like instances [6], later extended with fractal-dimension analyses of incidence graphs (Ansótegui et al. 2014, as tabulated in [4]); Newsham et al. [7] reported encouraging runtime regressions on community features. Zulkoski et al. [4] scaled the design to ∼7000 instances from SAT Competitions 2007–2016 with MapleCOMSPS runtimes and OLS/ridge regression over standardized features with full pairwise interactions. Their findings anchor our baseline expectations: no single parameter is significantly predictive in the aggregate; best combos reach adjusted $R^2 \approx 0.31$ (application), 0.56 (crafted), 0.66 (random); the size-plus-treewidth feature set carries almost no explanatory power on application instances (adjusted $R^2 \approx 0.05$); sub-category correlations are strong but *sign-flipping* (e.g. a mergeability ratio: Spearman +0.94 on argumentation, −0.73 on hardware); and pre-simplification plus corpus grouping materially change R^2 — a sensitivity we inherit and control by reporting both calibers. On the behavioral side, Zulkoski et al. show CDCL branching (VSIDS, LRB) is spatially local with respect to variable-incidence-graph communities (Gini 0.50–0.52 vs 0.16 for random branching) and temporally local (45–55% of decisions land in the most recent 1% of communities). Li et al. [3] document the hierarchical community structure of practical instances.

Locality and heterogeneity: theory. Bläsius et al. [8] prove that heterogeneity alone does not make random k-SAT easy (superpolynomial resolution size), whereas geometric locality yields small unsatisfiable subformulas findable in polynomial time — the strongest existing theorem-level support for "locality ⇒ tractability", and the regime map our generator navigates. Giráldez-Cru and Levy's locality-based generator variants [9] showed empirically that CDCL solvers exploit exactly this locality. Mull et al. [5] then proved average-case hardness for community-attachment instances with few communities, and showed experimentally that *community size, not formula size, dominates runtime*: instances differing 55× in variables but matched in community size differ ∼3× in runtime — an observational precursor of our interventional question. Jia, Moore and Strain [10] analyze planted SAT via solution hiding; their planting threshold is consistent with naive planted ensembles being CDCL-trivial, a pilot finding that motivated our non-planted difficulty carrier.

Proof complexity: the width backbone. Atserias and Dalmau [11] characterize resolution width combinatorially, yielding $w(F \vdash \bot) \geq \mathrm{tw}(\mathrm{primal}) - 2$ for 3-CNFs. Ben-Sasson and Wigderson [12] relate width to size, $S \geq \exp(\Omega((w - w_0)^2/n))$: a "width barrier" at $\approx\sqrt{n \log n}$. Atserias–Fichte–Thurley [14] and Beame et al. [15] connect CDCL to (bounded-width) resolution. Together: low treewidth ⇒ short proofs exist ⇒ CDCL with restarts finds them in polynomial time — the theoretical face of our E1 result (width-controlled instances are flat and trivial at all k). Correlation decay [16] and the overlap-gap property (see [17]) delineate where local information provably suffices or fails; we use these as interpretive scaffolding only, since practical CDCL carries components outside these models.

3 Problem Setup and Instance Generators

Locality kernel (`geo_random`). Fix $n = 400$ variables placed uniformly at random on a continuous 2-D unit torus. Each of $m = \alpha n$ clauses: pick the position of a uniformly chosen variable as the center, collect the variables within torus distance r (up to 1000 attempts to find a ball containing at least three variables), sample 3 distinct variables from that ball, and draw polarities uniformly (`planted=False`). Small r gives high intra-cluster variable reuse and sparse inter-cluster coupling; $r \to \infty$ recovers approximately uniform random 3-SAT. Because the torus diameter is $\sqrt{2}/2 \approx 0.71$, the r=0.8 and r=1.5 levels (whose r-ball spans all n variables) generate identical clause sets; sensitivity analyses treat them as one level. A *planted* variant draws polarities from σ-satisfaction modes built on a global assignment σ: for the k positions a clause satisfies σ on, a cardinality is drawn uniformly in $1 \dots k$ and then a subset of that cardinality; a decoy assignment β instead constrains the fill of the hidden clause fraction (the $q_{\text{hidden}} < 1$ arm). Clause-content statistics are therefore r-invariant by construction: redundancy histograms are identical across r (verified in pilots; measured mean σ-redundancy ≈2.2).

Width-controlled family. k-tree skeletons give constructive treewidth $\leq k$; matched random families at the same (n, α) provide controls. Planted and unplanted arms separate "having a solution planted" from "being structured".

Structure-destruction intervention (E3b). A bipartite swap chain (MCMC edge swaps on the variable–clause incidence graph, rejecting swaps that duplicate a variable in a clause or create complementary-literal pairs) destroys spatial locality while preserving every literal's degree *exactly*. The swapped instance is not a draw from the `geo_random` distribution; per our preregistration we report it as a structure-destruction intervention, and record solver-status flips per pair explicitly.

Topology × content factorial. The topology channel is original vs. swapped incidence graph. The content channel's unambiguous operationalization on our unsatisfiable-arm carrier is non-trivial (a distribution-preserving polarity resample is the identity by construction there); we registered this tension openly as AMEND-3 and deferred the interaction test rather than shipping an ill-defined arm. The topology channel is fully executed (E3b).

Threshold matching. Difficulty near threshold is U-shaped in density, so comparing arms at fixed absolute α confounds r-effects with threshold proximity. A dedicated threshold study (S0, §5) estimates the operational threshold $\alpha_c(r)$ per radius; the main scan samples at $\alpha = \alpha_c(r) + \Delta$, $\Delta \in \{+0.2, +0.4, +0.8\}$, plus a post-freeze fourth offset Δ=+0.6 at 4 of the 9 radii (deviation E3b-Δ).

4 Measures, Solvers, and Statistical Protocol

Difficulty. Primary: $\log_{10}(\text{conflicts} + 1)$ under CaDiCaL, conflict budget 10^6 (main scan) / 10^7 (threshold study). CaDiCaL's budget mechanism counts decided conflicts precisely, giving machine-independent censoring. Secondary: wall-clock (same machine), Glucose cross-checks. A 300 s subprocess-isolated wall guard returns status WALLTIMEOUT for runs where conflicts are sparse but propagation explodes; these rows carry no conflict count, are excluded from conflict analyses and from the Kaplan–Meier estimation (their censored-mass trend is reported separately, §6); their directional bias is declared (AMEND-2): truncation falls disproportionately on the UNSAT side near

threshold, which *raises* resolved-SAT rates and pushes $\alpha_c(r)$ upward — an anti-conservative bias for our gradient hypothesis, bounded by spot-checks at 10^8 conflicts / 480 s on flagged radii (§10).

Structural measures. Treewidth is reported as a bracket: min-fill upper bound (validated against k-tree constructions) and degeneracy-based lower bound. Community structure (Leiden, fixed resolution), modularity, spectral gap, redundancy histograms; all metrics are computed deterministically on the generated instance and stored with it.

Statistics. Repeated-measures ANOVA with seed as subject (seeds shared across all $r \times \Delta$ cells), partial $\eta^2 = F \cdot df_1/(F \cdot df_1 + df_2)$; Jonckheere–Terpstra ordered-alternative test by Monte Carlo permutation; paired Wilcoxon signed-rank for E3b pairs; Tobit (type-I, right-censored) and Kaplan–Meier as censoring analyses; Imai et al. [13] causal mediation (product-of-coefficients estimator) with bootstrap percentile ACME confidence intervals; Benjamini–Hochberg FDR across the primary family. Decision rules were preregistered: H2 confirmed by FDR-corrected $p < 0.05$ with partial $\eta^2 \geq 0.14$ or $\geq 10\times$ median endpoint ratio; "density dominates" triggered by $\eta^2 < 0.02$ and $p > 0.05$.

Preregistration and audit. Hypotheses, censoring semantics, r levels, and decision rules were frozen before the main scan (2026-09-09; scan completed 2026-09-10). The first *public* timestamp of the frozen documents, the amendment records AMEND-1–4, and the audit trail is the Zenodo deposit of 2026-09-15; AMEND-5, dated 2026-09-18, first appeared with the deposit's fourth version — we disclose this freeze-to-publication gap rather than let "preregistered" imply an independent registry. Throughout, *unblinding* denotes the first examination of main-scan outcome statistics; the S0 thresholds that dose the scan are examined beforehand by design, and the audit and its one material fix predate unblinding. Amendments are numbered and reason-stated: AMEND-1 (S0 budget 10^7), AMEND-2 (wall guard, bias declaration, spot-check protocol), AMEND-3 (content-arm deferral). A prespecification-vs-implementation audit records every deviation found by item-by-item inspection — including one material bug (walltimeout rows initially counted as UNSAT in threshold estimation) fixed *before* the main scan fired, and the post-freeze fourth dosing offset $\Delta{=}{+}0.6$ at 4 of the 9 radii (deviation E3b-Δ).

5 Results I: The Threshold Curve $\alpha_c(r)$ (S0)

S0 resolves 832 runs (720 preregistered grid jobs + 112 grid-floor extension jobs) at 10^7-conflict budget, 16 seeds per cell. With only resolved runs contributing to SAT rates (AMEND-2 caliber) and linear interpolation of the 0.5-crossing (Table 1; Fig. 1):

r	0.06	0.08	0.11	0.15	0.22	0.3	0.5	0.8	1.5
$\alpha_c(r)$	$< 2.5^b$	3.35	3.82	4.20	4.23	4.40^a	4.70^a	4.60^a	4.60^a
censoring	2.1%	0%	0%	0%	0%	92.5%	98.8%	100%	100%

Table 1: Operational threshold vs. locality radius. ^{a}measurement-saturated radii: the crossing is unidentifiable; the table carries the pilot coarse value, flagged low-confidence, with spot-checks per AMEND-2. ^{b}the SAT rate never reaches 0.5 on the explored grid ($\alpha \in [2.5, 3.1]$, rate $\sim$0.31 at the grid floor declining to $\sim$0.13): the boundary at $r = 0.06$ is bracketed from above, not pinpointed.

Three observations. *(i)* The operational threshold is not r-invariant. At $r = 0.06$ the SAT rate is already only $\sim$0.31 at density 2.5 and declines gently through 3.1 (censoring $\leq 2.1\%$): the boundary

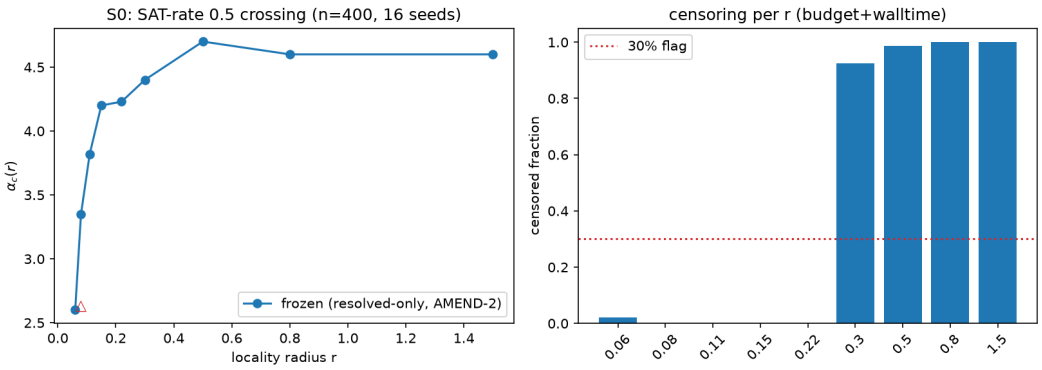


Figure 1: Left: the operational threshold $\alpha_c(r)$, resolved-only caliber (AMEND-2; △ marks the r=0.06 lattice-nearest anchor, the one radius whose grid has no 0.5-crossing). Right: censoring fraction per radius (budget + walltime) — the measurability boundary itself tracks r.

is bracketed below density 2.5. By $r = 0.22$ it has risen past 4.2 — a displacement of *at least* 1.7 density units. Locality does not merely re-scale difficulty around a fixed threshold, it *moves the threshold itself*. *(ii)* The curve saturates near the classical random-3-SAT threshold region [18] ($\alpha_c \to 4.267$ as $n \to \infty$), consistent with $r \to \infty$ recovering uniform random 3-SAT and with the width-backbone picture: small radii admit satisfiability far below it (the geometric small-subformula regime of [8]). *(iii)* At $r \geq 0.3$, near-threshold instances collapse out of measurable range (313 walltimeouts concentrated there): the question "where is the threshold" stops being answerable at fixed budget — a difficulty-regime change, bounded by spot-checks (§10).

6 Results II: Dose–Response at Matched Threshold Offset (E3a)

The main scan places 9 radii × 3 offsets above each radius's own $\alpha_c(r)$, 30 paired seeds, two solvers: 1740 runs, of which 1350 resolved and 390 censored (the per-arm grid appears in the journal version's appendix). The store additionally carries the 240 swap-arm rows of §7 (86 resolved, 154 censored), for 1980 runs with zero errored rows and 544 censored overall (27.5%, concentrated at $r \geq 0.3$); of those 544, 535 hit the 10^6-conflict budget and 9 the 300 s wall guard, which excludes at scale only in S0 (317 rows, 313 at $r \geq 0.3$). Figure 2 visualizes both arms.

Primary tests (CaDiCaL, Δ-pooled, seeds as subjects). RM-ANOVA on the Δ-aggregated complete-seed subset ($\Delta \in \{0.2, 0.4, 0.6, 0.8\}$, the 0.6 arm present at 4 of the 9 radii; $n = 28$ seeds observed at all radii): $F(8, 216) = 2808.5$, $p \approx 1.1 \times 10^{-213}$. *Scoring convention:* cell means average *decided* rows only (budget-censored and walltimeout runs are excluded, per the declared semantics), computed per (seed, radius) across the Δ arms present, so at heavily censored radii the mean rests on whichever arms resolved (unequal arm weights); a seed enters the subset only if every radius contributes at least one decided run (28 of 30 seeds qualify). Jonckheere–Terpstra (Monte Carlo permutation, 10^4 draws): $p \leq 10^{-4}$ — the resolution floor of the permutation distribution — for every (Δ, solver) cell; BH-FDR $q \leq 10^{-4}$ across the seven-member primary family. A censoring-robust variant (censored runs scored at the budget value) reproduces $p \leq 10^{-4}$, so the trend is not an artifact of dropping censored rows. Kaplan–Meier curves of time-to-decision (Fig. 3) make the same point visually: the median conflict count at decision rises from 23 ($r = 0.06$; identical under the Kaplan–Meier estimator) through 20,695 ($r = 0.22$, Kaplan–Meier median) to censoring

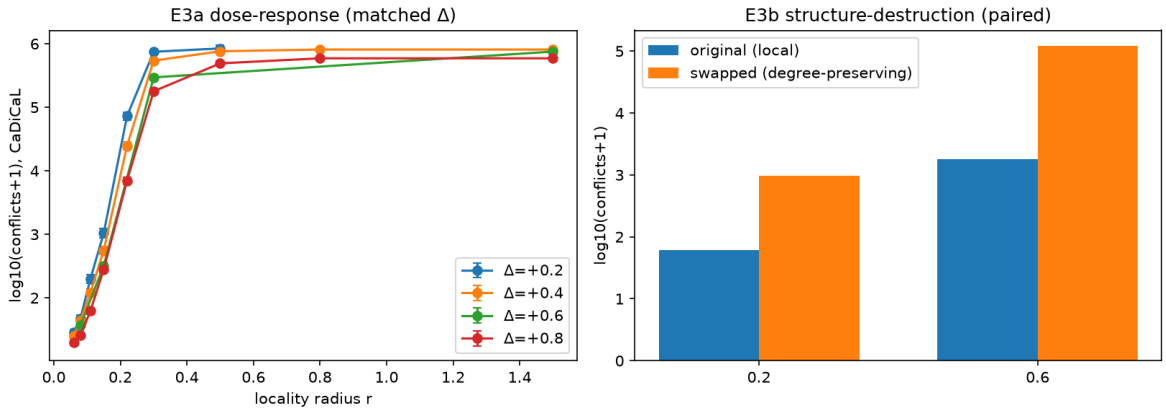


Figure 2: Left: dose–response of difficulty to the locality radius at matched threshold offsets (mean ± s.e.m., CaDiCaL; Δ=+0.8 endpoints 1.29 at r=0.06 and 5.77 at r=0.8). Right: paired structure-destruction effect (E3b).

plateaus — half or more undecided at $r \geq 0.5$, about a quarter at $r = 0.3$.

r	0.06	0.08	0.11	0.15	0.22	0.3	0.5
mean $\log_{10}$ conflicts	1.45	1.68	2.30	3.02	4.86	5.87[a]	5.93[a]
resolved n	30	30	30	30	30	9/30	2/30

Table 2: Dose–response at $\Delta = +0.2$ (the near-threshold arm, CaDiCaL). [a]means over surviving (easier) instances only; the censoring-aware Jonckheere–Terpstra bounds the underlying trend.

From $r = 0.06$ to $r = 0.22$ at the *same* threshold offset, difficulty climbs from ∼28 to ∼72,000 conflicts (Table 2) — **3.4 orders of magnitude** among decided instances, and ≥ 4.5 orders when the censored runs at $r \geq 0.3$ are counted at their budget lower bound. The $\Delta = +0.8$ arm repeats and extends the pattern (1.29 at r=0.06 → 5.77 at r=0.8, ≈4.5 orders). Partial η^2 at computable cells is 0.985–0.988 against the preregistered confirmation threshold of 0.14, and the $\Delta = +0.8$ arm's endpoint ratio $\sim 10^{4.5}$ dwarfs the preregistered 10× line. **H2 is confirmed under its own preregistered decision rule** (the statistical branch; the degree-*shape* premise holds in dispersion but not in the max-degree tail, ∼29% against the frozen <5% criterion — §10 quantifies the deviation and why the attribution does not rest on it). Glucose replicates the pattern.

Δ semantics under $\alpha_c(r)$ bias. For flagged radii the α_c carries a bounded (and for $r \geq 0.5$, unbounded) overestimate (§10); this shifts the precise horizontal placement of those cells but leaves the matched-comparison structure (within-r across Δ, within-Δ across r) intact, which is where the gradient evidence lives.

7 Results III: The Structure-Destruction Intervention (E3b)

Degree-preserving swapping holds every literal's occurrence count *exactly* fixed while destroying spatial locality. Paired Wilcoxon signed-rank:

As a sensitivity check, the effect is not an artifact of comparing across decision boundaries: within the 29 status-stable pairs (9 sat–sat at Δ=+0.2, 20 unsat–unsat at Δ=+0.6), swapped

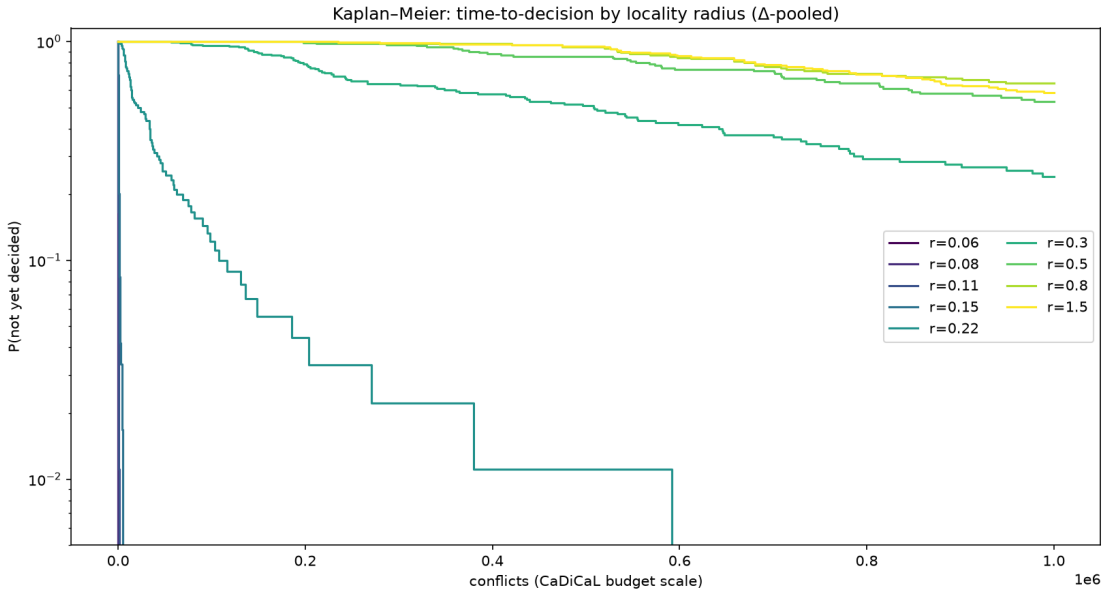


Figure 3: Kaplan–Meier estimates of time-to-decision (conflicts; CaDiCaL, Δ-pooled). Curves for small r drop to zero within the budget; curves for $r \geq 0.5$ plateau at high "not yet decided" probability (about a quarter at $r = 0.3$) — the measurability boundary of §5, seen at the solver level.

Δ	pairs	original mean logc	swapped mean logc	ratio	p
+0.2	33	1.79	2.99	×16	7.0×10^{-10}
+0.6	49	3.25	5.08	×68	6.0×10^{-13}

Table 3: E3b paired results. 53 of the 82 resolved pairs flip satisfiability status under the swap.

instances still take more conflicts in 25 of 29 (×11, $p = 3.9 \times 10^{-3}$ at +0.2; ×1.8, $p = 3.2 \times 10^{-4}$ at +0.6; pooled $p = 6.3 \times 10^{-7}$), with magnitudes below the pooled ratios of Table 3 because the flip stratum concentrates the extreme cases. Each arm attempts 120 swap pairings (4 radii × 30 seeds; original-arm counterparts come from the main-scan store); completions fall to 33 and 49 because the large-radius cells saturate the budget in either arm (at +0.2 the original arm decides 9/30 and 0/30 at r=0.3 and r=1.5, the swap arm 0/30 at both and 3/30 at r=0.15).

Together with E1, whose width-controlled instances are flat and trivial at all k (per-k medians pooling both sizes, unplanted arm: 7–10 conflicts across $k \in \{3, 5, 8, 12, 16, 20\}$, at $n \in \{300, 600\}$), the asymmetry is sharp: what CDCL difficulty responds to is *local geometry*, not any global topological summary.

8 Results IV: Mediation and the Feature Ladder (H1)

Imai et al. [13] causal mediation of the $r \to$ difficulty path (bootstrap 95% CI; total effect 3.26 log units): modularity ACME = 3.06 $[2.84, 3.37]$ (∼94% of total); spectral gap ACME = 3.92 $[3.68, 4.28]$ (suppression: direct effect −0.66 — consistent with a secondary channel, though marginal single-mediator fits can also produce this pattern under misspecification, so no claim rests on it); mean degree 2.34 $[2.15, 2.60]$; clustering 2.19 $[2.02, 2.44]$. All CIs exclude zero; the mediator set was

adapted after the outcome data was collected (AMEND-5, at an audit prior to unblinding), so the block is read as exploratory, supportive mechanism evidence rather than a confirmatory test, and the confirmatory weight rests with E3b.

The H1 feature ladder: on the main-scan data, ridge-CV R^2 rises from 0.610 (size features) and 0.664 (size+density) to **0.976** with structural features: on controlled families, structure carries predictive power that observation on random families cannot see (random-family baseline: 0.90 with zero structural gain). The preregistered nested OLS F-test on the same subset confirms the ladder: $F(60, 1283) = 297.8$, $p < 10^{-4}$ for the structural block beyond size+density. A gradient-boosted predictor with leave-one-r-out extrapolation (training never sees the target radius) reaches $R^2 = 0.76$: structural features generalize to unseen locality levels.

9 Results V: Placement on Real Instances (E4-prime)

All 59 fetched real instances (42 combinatorial/crafted from a public benchmark suite; 17 DIMACS classics) were solved under the S0 budget protocol: 22/59 decided. The DIMACS classics anchor the easy end: 15 SAT and 2 UNSAT, all decided within milliseconds (an audit caught a parsing artifact in the 1996 legacy DIMACS dialect, silently truncating these files; they were re-solved after the fix, with the full trail in the repository log). The combinatorial challenge instances mostly exceed the wall guard (hard anchor); file-name labels match our determinations on every checkable Heule instance, though one legacy `-yes` file is in fact UNSAT. Our generator's difficulty spectrum spans the middle of the real spectrum from textbook-easy to competition-hard. The industrial-diversity caveat stands (§10).

10 Threats to Validity

Construct. "Locality radius r" is one spatial operationalization on a 2-D torus host; other hosts and other notions may disagree. Our claim is about r as defined. Difficulty is conflict counts under CDCL — a proxy for, not a definition of, hardness.

Internal. Five controls close the confounding routes we identified: the variable count n; the density α (fixed or Δ-matched); the literal degree profile (degree-cap rejection in the planted arms; *exact* preservation under the E3b swap, the decisive control); the clause content (standardized polarity sampling, r-invariant redundancy); and the solver engine (two engines, both VSIDS-family CDCL; the registered local-search prediction of §11 is the declared probe beyond that family). Metrics are computed deterministically before solving, so the metric-computation channel cannot pseudo-correlate with runtime. Censoring is the residual internal threat: the wall guard is directionally biased (AMEND-2), and the direction was declared before unblinding. Spot-checks at 10^8 conflicts bound the α_c bias. At $r = 0.3$, SAT instances appear 0.2 below the operational threshold (bias $\in [-0.2, 0]$ density units); at $r \geq 0.5$, the ± 0.2 offsets fail to bracket the boundary at $10^8/480\,\mathrm{s}$, since every instance at $\alpha_c + 0.2$ is decided UNSAT while those at $\alpha_c - 0.2$ stay undecided. The measurability boundary is thus robust, and the flagged α_c values carry an unbounded-bias caveat. One material implementation deviation was caught by the prespecification audit *before* unblinding and fixed; the trail is public.

Literal-degree profile. Because density is Δ-matched to each radius's own threshold, the mean literal degree (3α: 8.4 at $r = 0.06$ vs. 14.7 at $r = 0.5$ in the Δ=+0.2 arms) necessarily co-varies with r; this is inherent to threshold matching, and standardized by measuring density as a distance from

each radius's own boundary. The *shape* of the literal-degree distribution, however, is monitored rather than enforced in the main `geo_random` family. Across the Δ=+0.2 cells (30 seeds each) the per-cell worst-case maximum degree spans 24–31 across r ($\sim$29% above the low-r endpoint), exceeding the <5% matching criterion laid down in the frozen H2 premise; the per-cell standard deviation stays within it (3.75 $\rightarrow$ 3.85, $\sim$3%). We disclose this deviation rather than silently incur it. It does not carry the causal attribution, which rests on the E3b swap alone: that intervention preserves the literal occurrence counts of every instance *exactly*, so the 16–68$\times$ paired difficulty shift cannot be carried by degrees. We do not lean on the direction of the residual shape drift, because the two tail readings point opposite ways: absolute tails grow toward the *large-r* end (where heavier tails are typically the *easier* regime for CDCL), while the relative dispersion falls (coefficient of variation $\sim$0.45 $\rightarrow$ 0.26 across r), i.e. the small-r end is the relatively heavier-tailed one. Either reading leaves the causal claim unchanged. A degree-capped `geo_random` variant remains an open robustness arm.

External. `geo_random` is an intervention vehicle, not a survey of practice; the real-instance placement is restricted to decided instances, with the industrial-diversity caveat documented (the canonical SAT 2024 main-track files were unreachable from our network; the deterministic preselection table (400 instances with hashes and 15-solver reference runtimes) is retained in the authors' working archive for future completion). All conclusions are scoped to CDCL-family solvers, direct CNF encodings, $n = 400$, and conflict-count difficulty.

Statistical conclusion. Paired designs with 30 seeds per cell, preregistered effect-size thresholds, FDR over the primary family, and censoring-robust cross-checks. No formal power analysis was performed for the negative branch: 30 seeds per cell was set by the compute budget at design time, and the prespecification audit's design-stage power note (a minimum detectable paired effect of $dz \geq 0.52$ at power 0.80 under pilot dispersion, i.e., $\sim$0.2–0.3 orders of magnitude between adjacent cells) does not extend to a power calculation under the realized censoring — the negative-branch trigger ($\eta^2 < 0.02$ with $p > 0.05$) therefore carries none, the confirmatory branch is insensitive to this (observed effects exceed the thresholds by orders of magnitude), but the negative branch's interpretability would have rested on it.

11 Discussion: What This Does and Does Not Say About P vs NP

Nothing we measure bears on P versus NP. The conjecture we test descends from a philosophical *intuition* about why the computational world feels tractable — "causes act locally" — and intuitions of this kind deserve attention precisely because they can be operationalized and then shot down.

Phase transitions and a portable diagnostic. A classical observation, dating to Cheeseman, Kanefsky and Taylor's "Where the really hard problems are" [19], holds that combinatorial problems are hardest near structural phase transitions. Our results give that observation a *causal, solver-relative* reading for SAT/CDCL: the locality radius moves the transition itself and restructures difficulty within matched offsets. The framework (one structural knob, threshold-matched dosing, a structure-destruction intervention, censoring-aware statistics) is offered as a *portable diagnostic protocol*: any claim of the form "structural property P helps algorithm A on family F" can be subjected to the same treatment. We explicitly do *not* claim our conclusions extend to local-search solvers; the contrast is a registered prediction: if difficulty is a structure$\times$algorithm relation, a

stochastic local-search solver should respond to r differently, possibly inversely. We deliberately leave that test unexecuted here — a result produced after unblinding would be post-hoc by construction — so the prediction stays falsifiable and open; the generator, seeds, and the probSAT source ship with the snapshot, putting the test within reach of any reader.

Tractability is a relational property. E1 (width-controlled instances are flat) and the threshold curve together suggest that difficulty lives neither in global topology alone nor in density alone, but in the interaction between spatial organization and the search dynamics of a particular algorithm family. CDCL branching is measurably local [4]; our experiments ask whether the local geometry of the *instance* is what that locality feeds on. A "yes" makes the philosophical conjecture precise enough to be useful; a "no" makes it precise enough to be abandoned — either outcome is progress over an unfalsifiable slogan.

The measurability boundary is a finding. That near-threshold instances at $r \geq 0.3$ collapse beyond our budgets is not an inconvenience to be apologized for; it is the signature of a regime change. Between $r = 0.06$ (threshold bracketed cleanly at tiny budgets) and $r = 0.5$ (threshold invisible at 10^7 conflicts), the *character* of the difficulty phase transition changes — exactly what a locality-causal story predicts, and what a density-only story has no reason to produce.

Intervention as a community practice. Much of structure–difficulty research has leaned on observational regression — with the field's own flagship study concluding that no structural measure is strongly predictive of CDCL performance in the aggregate [4]. Intervention designs (one knob, everything else pinned, preregistered decision rules, censoring declared) are a standard remedy in empirical sciences that have faced the same impasse. The generator, the swap chain, the statistical protocol and the audit trail are released so that the next conjecture can be tested the same way.

12 Reproducibility Statement

AI-assisted tools were used in an assistive capacity for literature search and citation verification, implementing the generators, solvers, and statistical pipelines, executing the experiment matrix, and manuscript preparation (full disclosure in the accompanying Zenodo snapshot, DOI: 10.5281/zenodo.22777748).

All instances are deterministically generated from $(\text{family}, \text{params}, \text{seed})$ tuples; the result store keys rows by $(\text{family}, \text{params}, \text{seed}, \text{solver})$ with idempotent resume-per-row semantics; every figure regenerates from the SQLite stores by committed scripts; result databases and manifests are version-controlled. The full experiment matrix, budgets, the preregistration with numbered amendments, and the prespecification-vs-implementation audit are in the repository.

References

[1] J. Marques-Silva, I. Lynce, S. Malik. Conflict-driven clause learning SAT solvers. In *Handbook of Satisfiability*, 2nd ed., FAIA 336 (2021). doi:10.3233/faia200987.

[2] SAT Competition series (2002–2026). Official archives and results: `https://satcompetition.github.io/` (online resource).

[3] C. Li et al. On the hierarchical community structure of practical Boolean formulas. *SAT 2021*, LNCS 12831, pp. 359–376. doi:10.1007/978-3-030-80223-3_25; arXiv:2103.14992.

[4] E. Zulkoski, R. Martins, C. Wintersteiger, J. Liang, K. Czarnecki, V. Ganesh. The effect of structural measures and merges on SAT solver performance. *CP 2018*, LNCS 11008, pp. 436–452. doi:10.1007/978-3-319-98334-9_29.

[5] N. Mull, D. Fremont, S. A. Seshia. On the hardness of SAT with community structure. *SAT 2016*, LNCS 9710, pp. 141–159. doi:10.1007/978-3-319-40970-2_10; arXiv:1602.08620.

[6] C. Ansótegui, M. L. Bonet, J. Levy. Towards industrial-like random SAT instances. *IJCAI 2009*. Companion: *On the structure of industrial SAT instances*, *CP 2009*, pp. 127–141, doi:10.1007/978-3-642-04244-7_13.

[7] Z. Newsham, V. Ganesh, S. Fischmeister, G. Audemard, L. Simon. Impact of community structure on SAT solver performance. *SAT 2014*, LNCS 8561, pp. 252–268. doi:10.1007/978-3-319-09284-3_20.

[8] T. Bläsius, T. Friedrich, A. Göbel, J. Levy, R. Rothenberger. The impact of heterogeneity and geometry on the proof complexity of random satisfiability. *SODA 2021*, pp. 42–53; *Random Structures & Algorithms* 63(4):885–941 (2023). doi:10.1002/rsa.21168.

[9] J. Giráldez-Cru, J. Levy. Locality in random SAT instances. *IJCAI 2017*, pp. 638–644. doi:10.24963/ijcai.2017/89.

[10] H. Jia, C. Moore, D. Strain. Generating hard satisfiable formulas by hiding solutions deceptively. *AAAI-05*, pp. 384–389.

[11] A. Atserias, V. Dalmau. A combinatorial characterization of resolution width. *J. Computer and System Sciences* 74(3):323–334. doi:10.1016/j.jcss.2007.06.025.

[12] E. Ben-Sasson, A. Wigderson. Short proofs are narrow — resolution made simple. *J. ACM* 48(2):149–169 (2001). doi:10.1145/375827.375835.

[13] K. Imai, L. Keele, D. Tingley, T. Yamamoto. Unpacking the black box of causality: Learning about causal mechanisms from experimental and observational studies. *American Political Science Review* 105(4):765–789 (2011). doi:10.1017/S0003055411000414.

[14] A. Atserias, J. K. Fichte, M. Thurley. Clause-learning algorithms with many restarts and bounded-width resolution. *JAIR* 40 (2011), pp. 353–373. doi:10.1613/jair.3152.

[15] P. Beame, H. Kautz, A. Sabharwal. Understanding the power of clause learning. *IJCAI 2003*; extended version *Towards understanding and harnessing the potential of clause learning*, *JAIR* 22:319–351 (2004). doi:10.1613/jair.1410.

[16] D. Weitz. Counting independent sets up to the tree threshold. *STOC 2006*, pp. 140–149. doi:10.1145/1132516.1132538.

[17] D. Gamarnik. The overlap gap property: a topological barrier to optimizing over random structures. *PNAS* 118(41) (2021). doi:10.1073/pnas.2108492118.

[18] M. Mézard, G. Parisi, R. Zecchina. Analytic and algorithmic solution of random satisfiability problems. *Science* 297(5582):812–815 (2002). doi:10.1126/science.1073287.

[19] P. Cheeseman, B. Kanefsky, W. Taylor. Where the really hard problems are. *IJCAI 1991*, pp. 331–337.

Reproducibility quick start.

```
mkdir cvh && cd cvh  # fetch + unpack the snapshot (DOI 10.5281/zenodo.22777748)
python3 -m venv .venv && .venv/bin/pip install -r requirements.txt
bash scripts/status.sh && .venv/bin/python experiments/e3a_analysis.py
.venv/bin/python experiments/e5_baseline.py results/m2_main.db results/m2_e1e2.db
```

All experiments completed 2026-09-10 (E4 DIMACS re-solved 2026-09-12 after a parsing fix); re-runs resume per row, skipping decided jobs. Both engines run through the python-sat 1.9.dev15 bindings (CaDiCaL 1.5.3, Glucose 4), stock configurations.